setup and generate steaming avater

1. Go to the server has a 4090 gpu
2. Install conda

```
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
bash Miniconda3-latest-Linux-x86_64.sh
source ~/.bashrc
```

3. Create and activate conda virtual env

```
conda create -n flashhead python=3.10
conda activate flashhead
```

4. Install pytorch

```
pip install torch==2.7.1 torchvision==0.22.1 --index-url
https://download.pytorch.org/whl/cu128
```

5. Download soulx-flashhead_1.3B

```
git clone https://github.com/Soul-AILab/SoulX-FlashHead.git
cd SoulX-FlashHead
```

6. Install dependencies

```
pip install -r requirements.txt
pip install ninja
pip install flash_attn==2.8.0.post2 --no-build-isolation
pip install sageattention==1.0.6 --no-build-isolation
```

7. Install FFmpeg

```
conda install -c conda-forge ffmpeg==7
```

8. install model

```
pip install "huggingface_hub[cli]"
hf download Soul-AILab/SoulX-FlashHead-1_3B --local-dir ./models/SoulX-FlashHead-1_3B
hf download facebook/wav2vec2-base-960h --local-dir ./models/wav2vec2-base-960h
```

9. run script to generate stream video

```
# Infer with [Lite-Model] on single GPU
bash inference_script_single_gpu_lite.sh
# Real-time inference speed can be supported on single RTX-4090 (up to
3 concurrent).
```